

# Assignment 3 - Numerical Linear Algebra

Arthur Rabello Oliveira<sup>1</sup>

27/05/2025

## Abstract

We design and test a function `to_hessemberg(A)` that reduces an arbitrary square matrix to (upper) Hessenberg form with Householder reflectors, returns the reflector vectors, the compact Hessenberg matrix  $H$ , and the accumulated orthogonal factor  $Q$ , verifying numerically that  $A = QH Q^*$  and  $Q^* Q = I$  for symmetric and nonsymmetric inputs of orders  $10 - 10000$ . Timings confirm the expected  $O(\text{something})$  cost and reveal the  $2 \times$  speed-up attainable for symmetric matrices through trivial bandwidth savings. Leveraging this routine, we investigate the spectral structure of orthogonal matrices: we show that all eigenvalues lie on the unit circle, analyse the consequences for the power method and inverse iteration, and obtain a closed-form spectrum for generic  $2 \times 2$  orthogonals. Random  $4 \times 4$  orthogonal matrices generated via QR factorisation are then reduced to Hessenberg form; the eigenvalues of their trailing  $2 \times 2$  blocks are computed analytically and reused as fixed shifts in the QR iteration, where experiments demonstrate markedly faster convergence. Throughout, every algorithm is documented and supported by commented plots that corroborate the theoretical claims.

## Contents

1. Introduction .....	2
2. Hessemberg Reduction (Problem 1) .....	3
2.1. Calculating the Householder Reflectors (a) .....	3
2.2. Evaluating the Function (b), (c), (d) .....	5
2.2.1. Complexity (c) .....	12
2.2.2. The Symmetric Case (d) .....	12
3. Orthogonal Matrices (Problem 2) .....	13
3.1. Eigenvalues and Iterative Methods (a) .....	13
3.2. The $2 \times 2$ Case (b) .....	13
3.3. Random Orthogonal Matrices (c) .....	13
3.4. Shift With an Eigenvalue (d) .....	13
4. Conclusion .....	13
Bibliography .....	13

---

<sup>1</sup>Escola de Matemática Aplicada, Fundação Getúlio Vargas (FGV/EMAp), email: [arthur.oliveira.1@fgv.edu.br](mailto:arthur.oliveira.1@fgv.edu.br)

# 1. Introduction

One could calculate the eigenvalues of a square matrix using the following algorithm:

1. Compute the  $n$ -th degree polynomial  $\det(A - \lambda I) = 0$ ,
2. Solve for  $\lambda$  (somehow).

On step 2, the eigenvalue problem would have been reduced to a polynomial root-finding problem, which is awful and extremely ill-conditioned. From the [previous assignment](#) we know that in the denominator of the relative condition number  $\kappa(x)$  there's a  $|x - n|$ . So  $\kappa(x) \rightarrow \infty$  when  $x \rightarrow 0$ . As an example, consider the polynomial

$$p(x) = (x - 2)^9 = x^9 - 18x^8 + 144x^7 - 672x^6 + 2016x^5 - 4032x^4 + 5376x^3 - 4608x^2 + 2304x - 512 \quad (1)$$

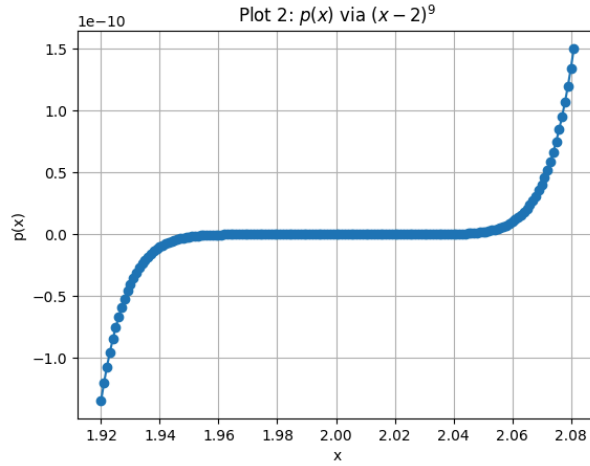


Figure 1:  $p(x)$  via the coefficients in [eq. \(1\)](#)




Figure 2:  $p(x)$  via  $(x - 2)^9$

[Figure 2](#) shows a smooth curve, while [Figure 1](#) shows a weird oscillation around  $x = 0$  (And pretty much everywhere else if the reader is sufficiently persistent).

This is due to the round-off errors when  $x \approx 0$  and the big coefficients of the polynomial. In general, polynomial are very sensitive to perturbations in the coefficients, which is why rootfinding is a bad idea to find eigenvalues.

Here we discuss aspects of some iterative eigenvalue algorithms, such as power iteration, inverse iteration, and QR iteration.

## 2. Hessemberg Reduction (Problem 1)

### 2.1. Calculating the Householder Reflectors (a)

The following function calculates the Householder reflectors that reduce a matrix to Hessenberg form. It returns the reflector vectors, the compact Hessenberg matrix  $H$ , and the accumulated orthogonal factor  $Q$ .

```
1  import numpy as np
2  import time
3  from typing import List, Tuple
4
5
6  def build_householder_unit_vector(
7      target_vector: np.ndarray
8  ) -> np.ndarray:
9
10     """
11     Builds a Householder unit vector
12
13     Args:
14         1. target_vector (np.ndarray): Column vector that we want to annihilate
15            (size  $\geq 1$ ).
16
17     Returns:
18         np.ndarray:
19             The normalised Householder vector ( $\|v\|_2 = 1$ ) with a real first
20             component.
21
22     Raises:
23         1. ValueError: If 'target_vector' has zero length.
24     """
25
26     if target_vector.size == 0:
27         raise ValueError("The target vector is empty; no reflector needed.")
28
29     vector_norm: float = np.linalg.norm(target_vector)
30
31     if vector_norm == 0.0: #nothing to annihilate – return canonical basis
32         vector
33
34     householder_vector: np.ndarray = np.zeros_like(target_vector)
35     householder_vector[0] = 1.0
36     return householder_vector
37
38     sign_correction: float = (
```

```

35         1.0 if target_vector[0].real >= 0.0 else -1.0
36     )
37     copy_of_target_vector: np.ndarray = target_vector.copy()
38     copy_of_target_vector[0] += sign_correction * vector_norm
39     householder_vector: np.ndarray = (
40         copy_of_target_vector / np.linalg.norm(copy_of_target_vector)
41     )
42     return householder_vector
43
44
45 def to_hessenberg(
46     original_matrix: np.ndarray,
47 ) -> Tuple[List[np.ndarray], np.ndarray, np.ndarray]:
48
49     """
50     Reduce 'original_matrix' to upper Hessenberg form by Householder
51     reflections.
52
53     Args
54
55         1. original_matrix (np.ndarray): Real or complex square matrix of order
56         'matrix_order'.
57
58     Returns
59
60         Tuple consisting of:
61
62         1. householder_reflectors_list (List[np.ndarray])
63         2. hessenberg_matrix (np.ndarray)
64         3. accumulated_orthogonal_matrix (np.ndarray) s.t.
65             original_matrix = Q · H · QH
66
67     Raises
68
69         1. ValueError: If 'original_matrix' is not square.
70
71     """
72
73     working_matrix: np.ndarray = np.asarray(original_matrix).copy()
74
75     if working_matrix.shape[0] != working_matrix.shape[1]:
76         raise ValueError("Input matrix must be square.")
77
78     matrix_order: int = working_matrix.shape[0]
79     accumulated_orthogonal_matrix: np.ndarray = np.eye(
80         matrix_order, dtype=working_matrix.dtype
81     )
82     householder_reflectors_list: List[np.ndarray] = []

```

```

78     for column_index in range(matrix_order - 2): #extract the part of column
        'column_index' that we want to zero out
79         target_column_segment: np.ndarray = working_matrix[
80             column_index + 1 :, column_index
81         ]
82
83         householder_vector: np.ndarray = build_householder_unit_vector(
84             target_column_segment
85         ) #build Householder vector for this segment
86         householder_reflectors_list.append(householder_vector)
87
88         #expand it to the full matrix dimension
89         expanded_householder_vector: np.ndarray = np.zeros(
90             matrix_order, dtype=working_matrix.dtype
91         )
92         expanded_householder_vector[column_index + 1 :] = householder_vector
93
94
95         working_matrix -= 2.0 * np.outer(
96             expanded_householder_vector,
97             expanded_householder_vector.conj().T @ working_matrix,
98         ) #apply reflector from BOTH sides
99         working_matrix -= 2.0 * np.outer(
100             working_matrix @ expanded_householder_vector,
101             expanded_householder_vector.conj().T,
102         )
103
104         #accumulate Q
105         accumulated_orthogonal_matrix -= 2.0 * np.outer(
106             accumulated_orthogonal_matrix @ expanded_householder_vector,
107             expanded_householder_vector.conj().T,
108         )
109
110     hessenberg_matrix: np.ndarray = working_matrix
111     return (
112         householder_reflectors_list,
113         hessenberg_matrix,
114         accumulated_orthogonal_matrix,
115     )

```

We will evaluate this function in [Section 2.2](#).

## 2.2. Evaluating the Function (b), (c), (d)

We present another algorithm for evaluating the function `to_hessenberg(A)` for random matrices of various sizes, inputed by the user, which also gets to choose if symmetric matrices will be generated or not.

```

1  import numpy as np
2  import time
3  import pandas as pd
4  import matplotlib.pyplot as plt
5  import math
6  from IPython.display import display, Markdown
7  from ast import literal_eval
8
9  #RANDOM MATRIX GENERATOR
10 def generate_random_matrix(n:int, distribution:str="normal",
11                             symmetric:bool=False, seed:int|None=None):
12     rng = np.random.default_rng(seed)
13     if distribution == "normal":
14         A = rng.standard_normal((n, n))
15     elif distribution == "uniform":
16         A = rng.uniform(-1.0, 1.0, size=(n, n))
17     else:
18         raise ValueError("distribution must be 'normal' or 'uniform'")
19     return (A + A.T) / 2.0 if symmetric else A
20
21
22 #REFLECTOR CALCULATOR
23 def _house_vec(x:np.ndarray) -> np.ndarray:
24
25     """
26     Builds a Householder reflector for a given column vector x.
27     Args:
28         x (np.ndarray): Column vector to be transformed.
29     Returns:
30         np.ndarray: Normalised Householder vector with a real first component.
31     Raises:
32         None
33     """
34
35     sigma = np.linalg.norm(x)
36     if sigma == 0.0:
37         e1 = np.zeros_like(x)
38         e1[0] = 1.0
39         return e1
40     sign = 1.0 if x[0].real >= 0.0 else -1.0
41     v = x.copy()
42     v[0] += sign * sigma
43     return v / np.linalg.norm(v)
44
45 def hessenberg_reduction(A_in:np.ndarray, symmetric:bool=False,
46                           accumulate_q:bool=True):

```

```

46
47     """
48     Reduces a matrix to upper Hessenberg form using Householder reflections.
49     Args:
50         A_in (np.ndarray): Input matrix to be reduced.
51         symmetric (bool): If True, treat the matrix as symmetric and reduce to
52                             tridiagonal form.
53         accumulate_q (bool): If True, accumulate the orthogonal matrix Q.
54     Returns:
55         Tuple[np.ndarray, np.ndarray]: The reduced matrix in Hessenberg form
56                                         and the orthogonal matrix Q.
57     Raises:
58         None
59     """
60
61     A = A_in.copy()
62     n = A.shape[0]
63     Q = np.eye(n, dtype=A.dtype)
64
65     if not symmetric:    #GENERAL caSe
66         for k in range(n-2):
67             v = _house_vec(A[k+1:, k])
68             w = np.zeros(n, dtype=A.dtype)
69             w[k+1:] = v
70             A -= 2.0 * np.outer(w, w.conj().T @ A)
71             A -= 2.0 * np.outer(A @ w, w.conj().T)
72             if accumulate_q:
73                 Q -= 2.0 * np.outer(Q @ w, w.conj().T)
74         return A, Q
75
76     #SYMMETRIC TRIDIAGONAL CASE
77     for k in range(n-2):
78         x = A[k+1:, k]
79         v = _house_vec(x)
80         beta = 2.0
81
82         w = A[k+1:, k+1:] @ v    #trailing submatrix rank-2 update ( $A \leftarrow A - v w^T - w v^T$ )
83         tau = beta * 0.5 * (v @ w)
84         w -= tau * v
85         A[k+1:, k+1:] -= beta * np.outer(v, w) + beta * np.outer(w, v)
86
87         new_val = -np.sign(x[0]) * np.linalg.norm(x)    #store the single sub-
88                                                         diagonal element, zero the rest
89         A[k+1, k] = new_val
90         A[k, k+1] = new_val

```

```

88         A[k+2:, k] = 0.0
89         A[k, k+2:] = 0.0
90
91         if accumulate_q: #accumulate Q if requested
92             Q[:, k+1:] -= beta * np.outer(Q[:, k+1:] @ v, v)
93
94     A = np.triu(A) + np.triu(A, 1).T #force symmetry
95     return A, Q
96
97
98 #VERIFYING PART
99 def verify_factorisation_once(n:int, dist:str, symmetric:bool, seed:int|None):
100
101     """
102     Verifies the factorisation of a random matrix of size n.
103     Args:
104         n (int): Size of the matrix.
105         dist (str): Distribution type ('normal' or 'uniform').
106         symmetric (bool): Whether the matrix is symmetric.
107         seed (int | None): Random seed for reproducibility.
108     Returns:
109         None
110     Raises:
111         None
112     """
113
114     A = generate_random_matrix(n, dist, symmetric, seed)
115     T, Q = hessenberg_reduction(A, symmetric=symmetric)
116     res_fact = np.linalg.norm(A - Q @ T @ Q.T)
117     res_orth = np.linalg.norm(Q.T @ Q - np.eye(n))
118     colour = "green" if res_fact < 1e-11 else "red"
119     typ = "symmetric" if symmetric else "general"
120     display(Markdown(
121         f"**{n}x{n} {typ}** \n"
122         f"<span style='color:{colour}'>||A - Q T QT|| = {res_fact:.2e}</span> \n"
123         f"||QTQ - I|| = {res_orth:.2e}"
124     ))
125
126
127 def benchmark_hessenberg(size_list, dist:str, mode:str, seed:int|None,
128     reps_small:int=5):
129
130     """
131     Benchmark the Hessenberg reduction for various matrix sizes and types.
132     Args:

```



```

132     size_list (list of int): List of matrix sizes to test.
133     dist (str): Distribution type ('normal' or 'uniform').
134     mode (str): Matrix type ('general', 'symmetric', or 'both').
135     seed (int | None): Random seed for reproducibility.
136     reps_small (int): Number of repetitions for small matrices.
137     Returns:
138         pd.DataFrame: DataFrame containing the benchmark results.
139     Raises:
140         None
141     """
142
143     records = []
144     for n in size_list:
145         for sym in ([False, True] if mode=="both" else [mode=="symmetric"]):
146             A = generate_random_matrix(n, dist, sym, seed)
147
148             t0 = time.perf_counter()
149             hessenberg_reduction(A, symmetric=sym, accumulate_q=False)
150             probe = time.perf_counter() - t0
151             reps = reps_small if probe*reps_small >= 1.0 else math.ceil(1.0 /
152                                     probe)
153
154             times = []
155             for _ in range(reps):
156                 start = time.perf_counter()
157                 hessenberg_reduction(A, symmetric=sym, accumulate_q=False)
158                 times.append(time.perf_counter() - start)
159
160             records.append(dict(size=n,
161                                type="symmetric" if sym else "general",
162                                reps=reps,
163                                avg=np.mean(times)))
164
165     df = pd.DataFrame(records)
166     display(df.style.format({"avg": "{:.3e}"}).hide(axis="index"))
167
168     plt.figure(figsize=(7,5))
169     mark = {"general": "o", "symmetric": "s"}
170     for label, sub in df.groupby("type"):
171         plt.loglog(sub["size"], sub["avg"], marker=mark[label], ls="--",
172                    label=label)
173         if len(sub) > 1:
174             a,b = np.polyfit(np.log10(sub["size"]), np.log10(sub["avg"]), 1)
175             plt.loglog(sub["size"], 10**(b+a*np.log10(sub["size"])),
176                        "--", label=f"{label} fit ~ $n^{a:.2f}$")
177     plt.xlabel("matrix size (log)")

```


```

176 plt.ylabel("runtime [s] (log)")
177 plt.title("Hessenberg (general) vs Tridiagonal (symmetric)")
178 plt.grid(True, which="both", ls=":")
179 plt.legend(); plt.tight_layout(); plt.show()
180 return df
181
182
183 #===INTERACTIVE PART=====
184 try:
185     raw = input("\nMatrix sizes (Python list) (e.g): [64,128,256,512,1024]: ")
186     sizes = literal_eval(raw) if raw.strip() else [64,128,256,512,1024]
187 except Exception:
188     print("Bad list -> using default.")
189     sizes = [64,128,256,512,1024]
190
191 dist = input("Distribution ('normal'/'uniform') [normal]: ").strip().lower()
192 or "normal"
193 mode_txt = input("Matrix type g=general, s=symmetric, b=both [g]: ")
194 or "g"
195 mode = "symmetric" if mode_txt=="s" else "both" if mode_txt=="b" else "general"
196 seed_txt = input("Random seed (None/int) [None]: ").strip()
197 seed_val = None if seed_txt.lower() in {"", "none"} else int(seed_txt)
198
199 for sym in ([False, True] if mode=="both" else [mode=="symmetric"]): #accuracy
200     on largest size
201     verify_factorisation_once(max(sizes), dist, sym, seed_val)
202
203 benchmark_hessenberg(sizes, dist, mode, seed_val) #timings

```

The reader should be aware that my poor Dell Inspiro 5590 has crashed precisely 5 times while i was writing this. The runtime was around 4 minutes for a matrix  $A \approx 10^3 \times 10^3$ .

An expected output is:

1	64×64 general	 Python
2	$\ A - Q T Q^T\  = 7.51e-14$	
3	$\ Q^T Q - I\  = 7.07e-15$	
4		
5	64×64 symmetric	
6	$\ A - Q T Q^T\  = 4.83e-14$	
7	$\ Q^T Q - I\  = 7.39e-15$	
8		
9	128×128 general	
10	$\ A - Q T Q^T\  = 1.84e-13$	
11	$\ Q^T Q - I\  = 1.26e-14$	
12		
13	128×128 symmetric	
14	$\ A - Q T Q^T\  = 1.14e-13$	

```

15 ||QTQ - I|| = 1.25e-14
16
17 256×256 general
18 ||A - Q T QT|| = 4.70e-13
19 ||QTQ - I|| = 2.28e-14
20
21 256×256 symmetric
22 ||A - Q T QT|| = 2.78e-13
23 ||QTQ - I|| = 2.25e-14
24
25 512×512 general
26 ||A - Q T QT|| = 1.16e-12
27 ||QTQ - I|| = 4.10e-14
28
29 512×512 symmetric
30 ||A - Q T QT|| = 7.10e-13
31 ||QTQ - I|| = 4.09e-14
32
33 1024×1024 general
34 ||A - Q T QT|| = 3.05e-12
35 ||QTQ - I|| = 7.57e-14
36
37 1024×1024 symmetric
38 ||A - Q T QT|| = 1.84e-12
39 ||QTQ - I|| = 7.64e-14

```

As  $n$  grows, we observe that the residuals also grow, but still in machine precision. The difference between the symmetric and nonsymmetric cases are more pronounced in larger matrices.




Figure 3: Runtime of the Hessenberg reduction for ordinary and symmetric matrices

### 2.2.1. Complexity (c)

Figure 3 shows the expected  $O(n^3)$  complexity for the general case and  $O(n^2)$  for the symmetric case. The latter is better discussed in Section 2.2.2.

To understand why the complexity is  $O(n^3)$  in the general case, we can look at the algorithm. The outer loop runs  $n - 2$  times, and inside it, we have two matrix-vector products and two outer products, which are all  $O(n^2)$ . Thus, the total complexity is  $O(n^3)$ .

### 2.2.2. The Symmetric Case (d)

On the symmetric case we know that reflectors will be applied in only one side of the matrix, since  $v^T A = A v^T$ . That is precisely what the function `generate_random_matrix` does. Which cuts complexity from the expected  $O(n^3)$  seen in the previous section to a  $O(n^2)^2$ .

<sup>2</sup>See page 194 of Trefethen & Bau's Numerical Linear Algebra book

### **3. Orthogonal Matrices (Problem 2)**

#### **3.1. Eigenvalues and Iterative Methods (a)**

#### **3.2. The $2 \times 2$ Case (b)**

#### **3.3. Random Orthogonal Matrices (c)**

#### **3.4. Shift With an Eigenvalue (d)**

### **4. Conclusion**

### **Bibliography**